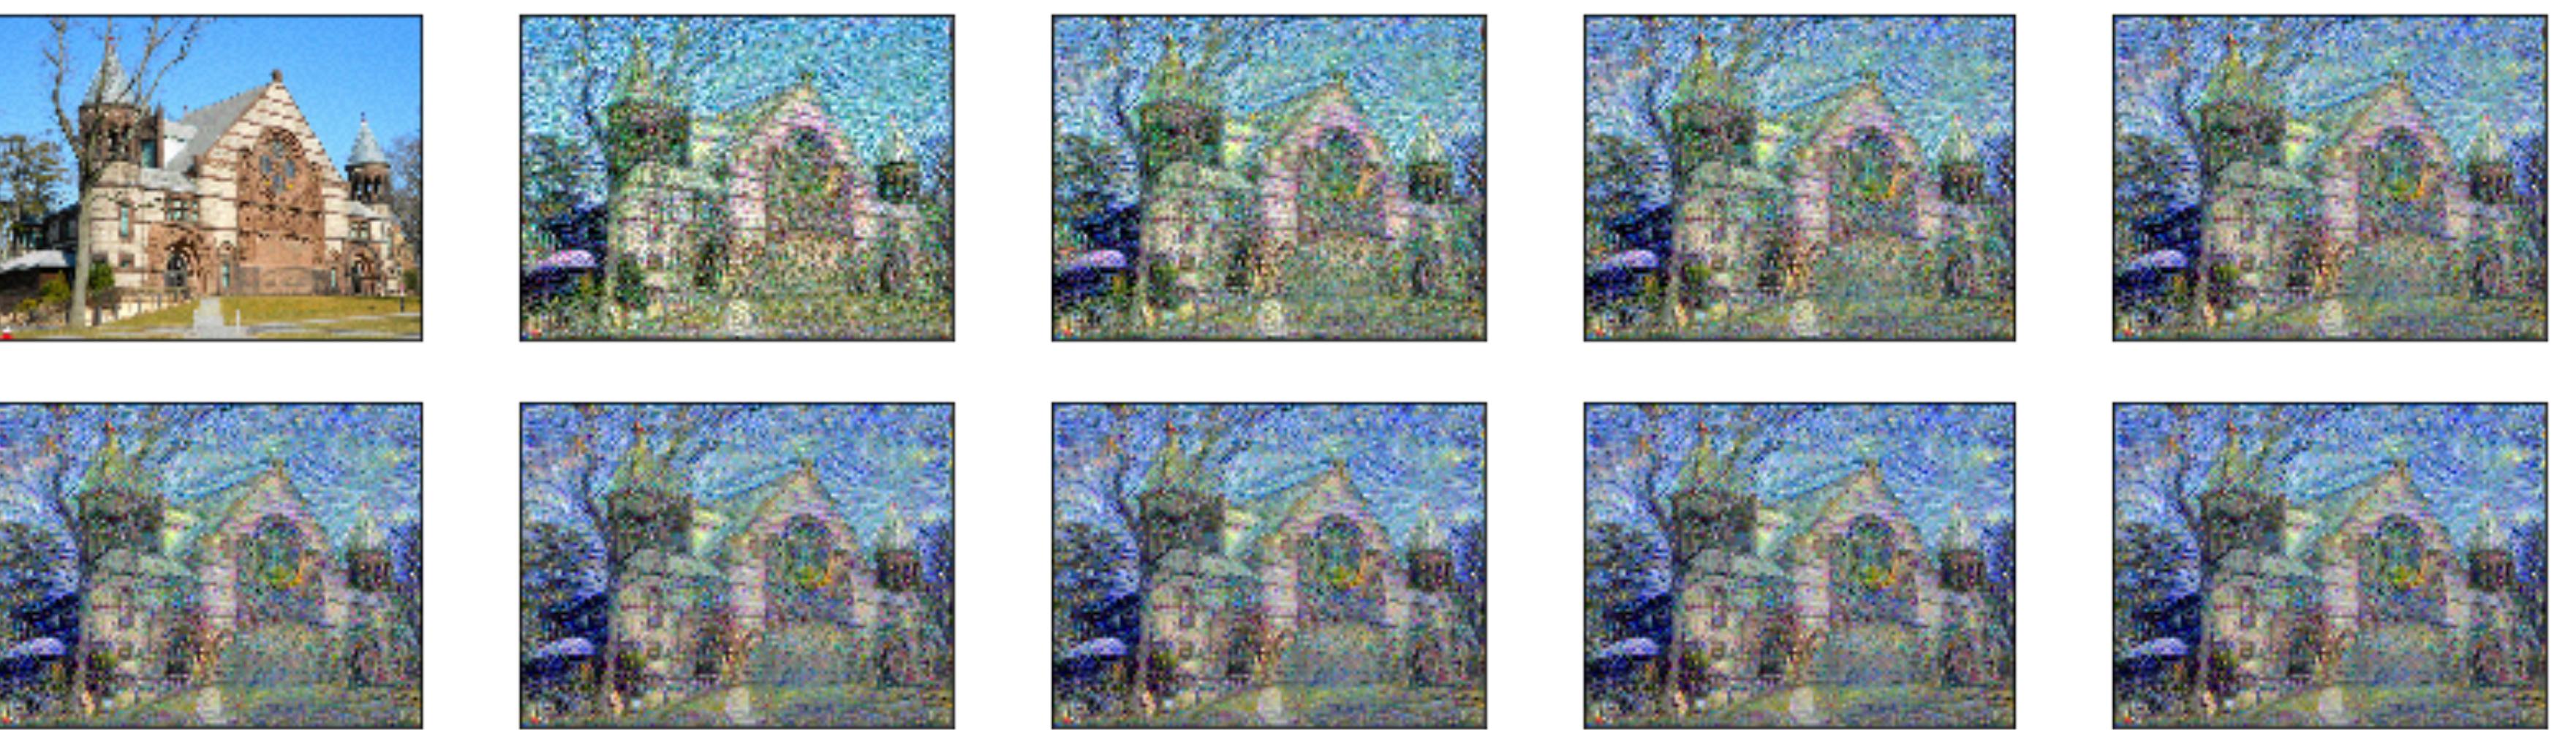


The Elements of Neural Artistic Style Transfer

Alexander Chen

alexanderKchen@princeton.edu

Princeton University, Department of Mathematics, Princeton, NJ



Motivation

- The fundamental problem of style transfer is to extract and apply the 'style' of artwork S, say a painting, music recording, or written work, to another artwork X to generate a new artwork Y that matches the content of X but with the style of S.
- The initial difficulty is how to define the 'style' of a picture or piece of literature, but in recent years deep neural networks have provided a way.
- Neural artistic style transfer (or NST for short) found its first success in using deep convolutional nets.
- A number of NST unsupervised learning methods are surveyed here including GANs and Gatys' original method.

Deep Methods for Artistic Transfer

Gatys' method

Let C be the content image and S the style image. We want to create x with the content of C and style of S.

We use a neural network F, the pre-trained VGG-19 CNN.

Suppose $F(x)$ represents the neuron weights of the network when x is fed in.

The content distance is

$$L_c(x, C) = \|F(x) - F(C)\|$$

The style distance is

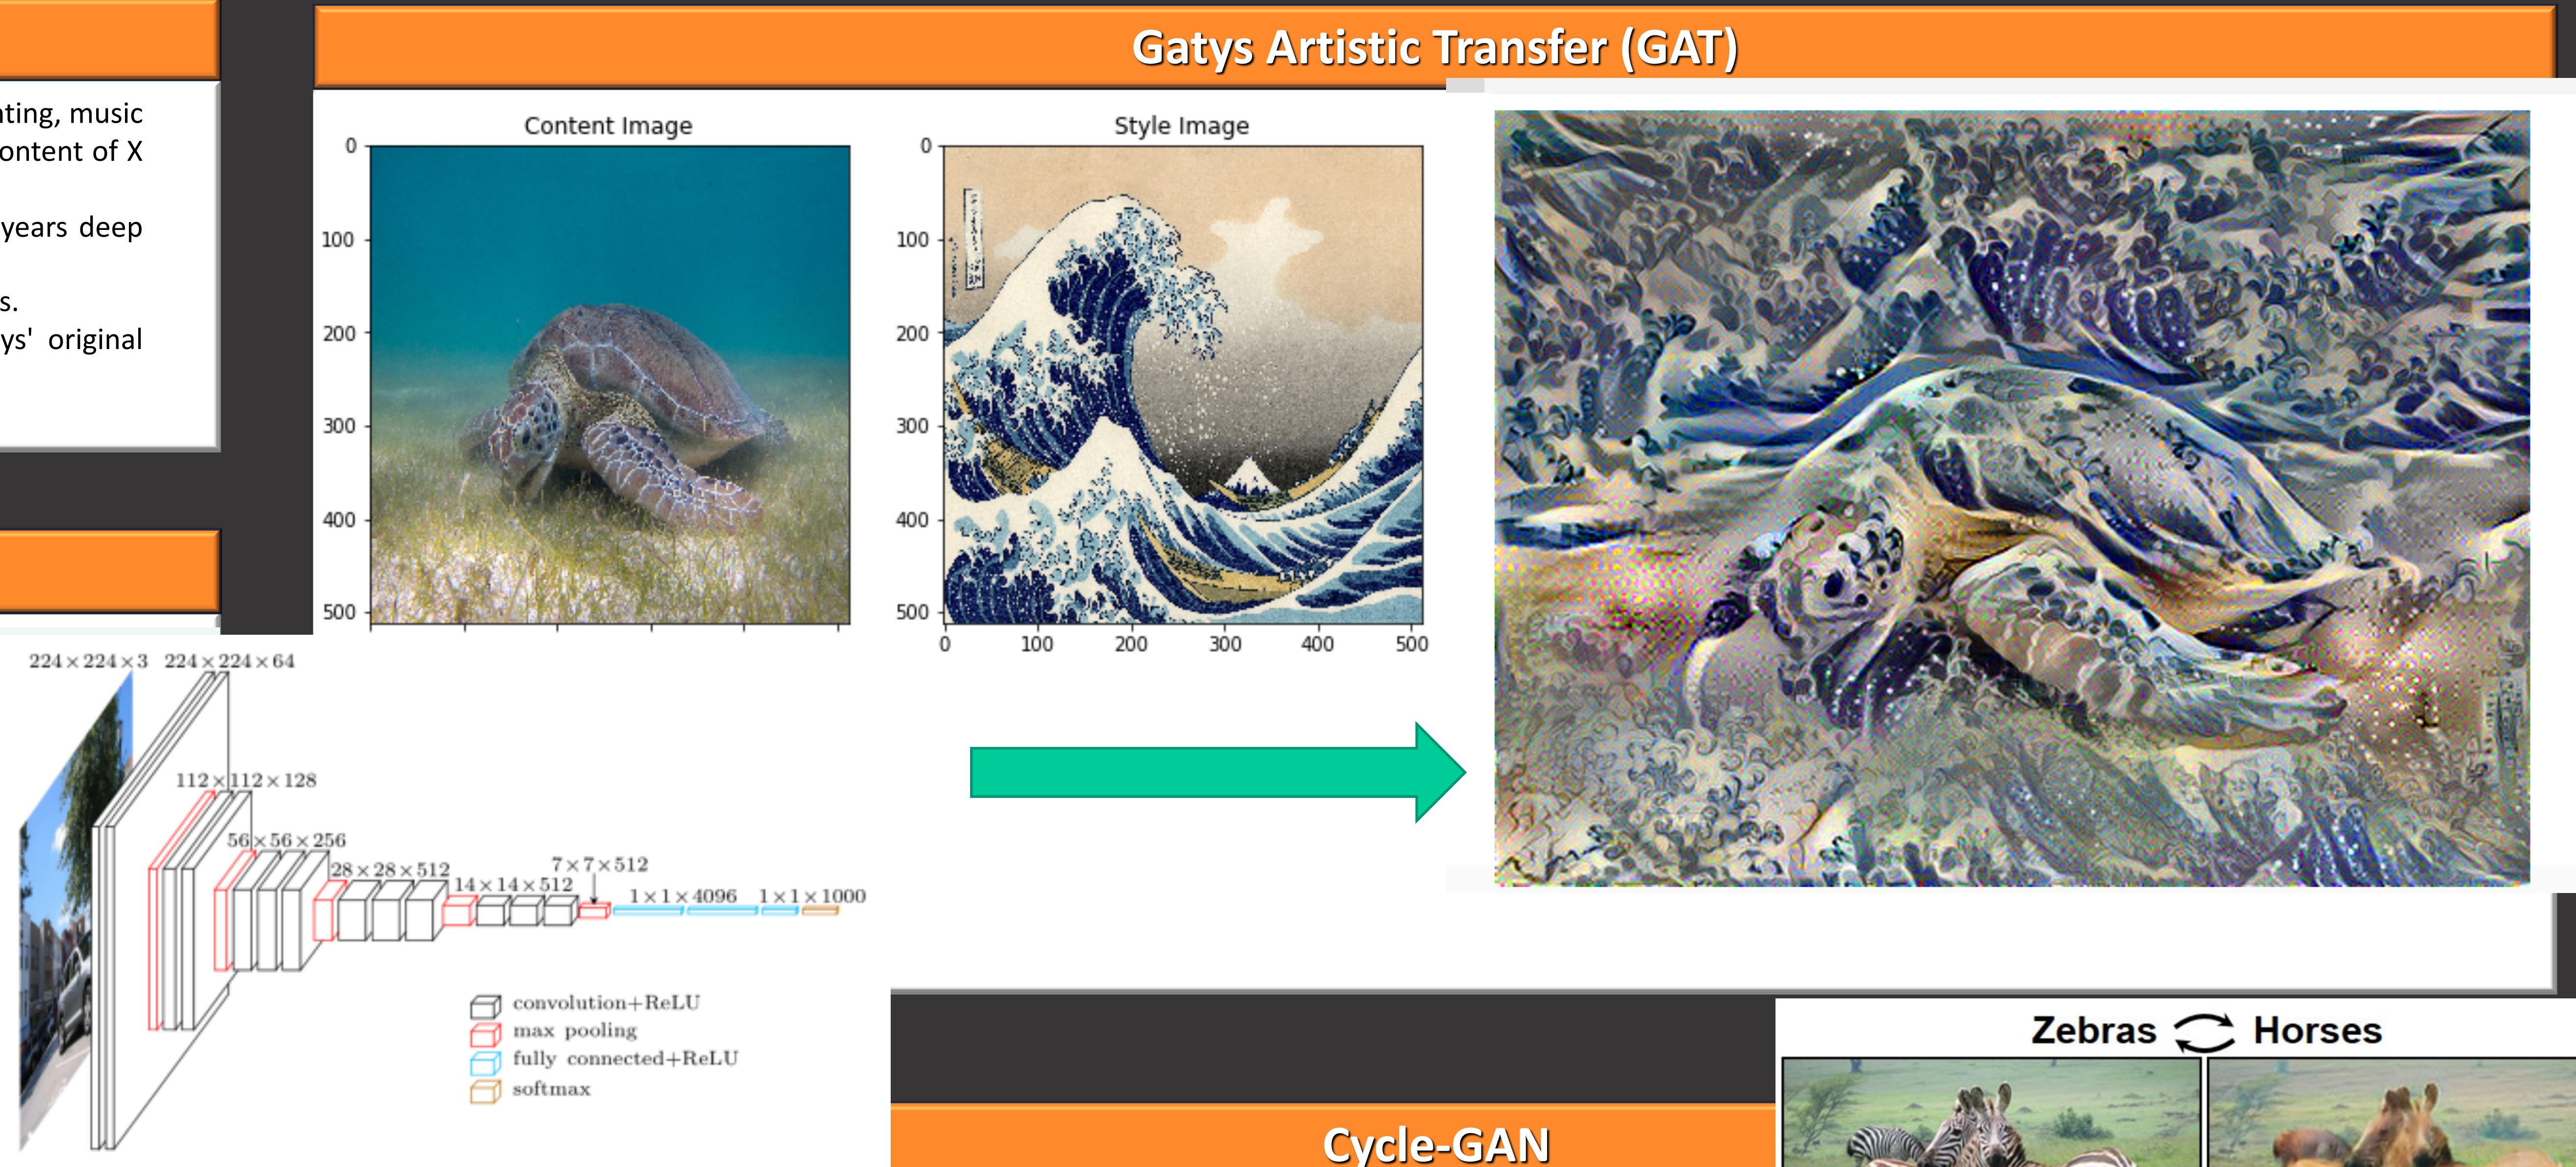
$$L_s(x, S) = \|G(F(x)) - G(F(S))\|$$

where G is the Gram matrix (of correlations). The goal is to make the given x as close as possible to the style of S and content of C

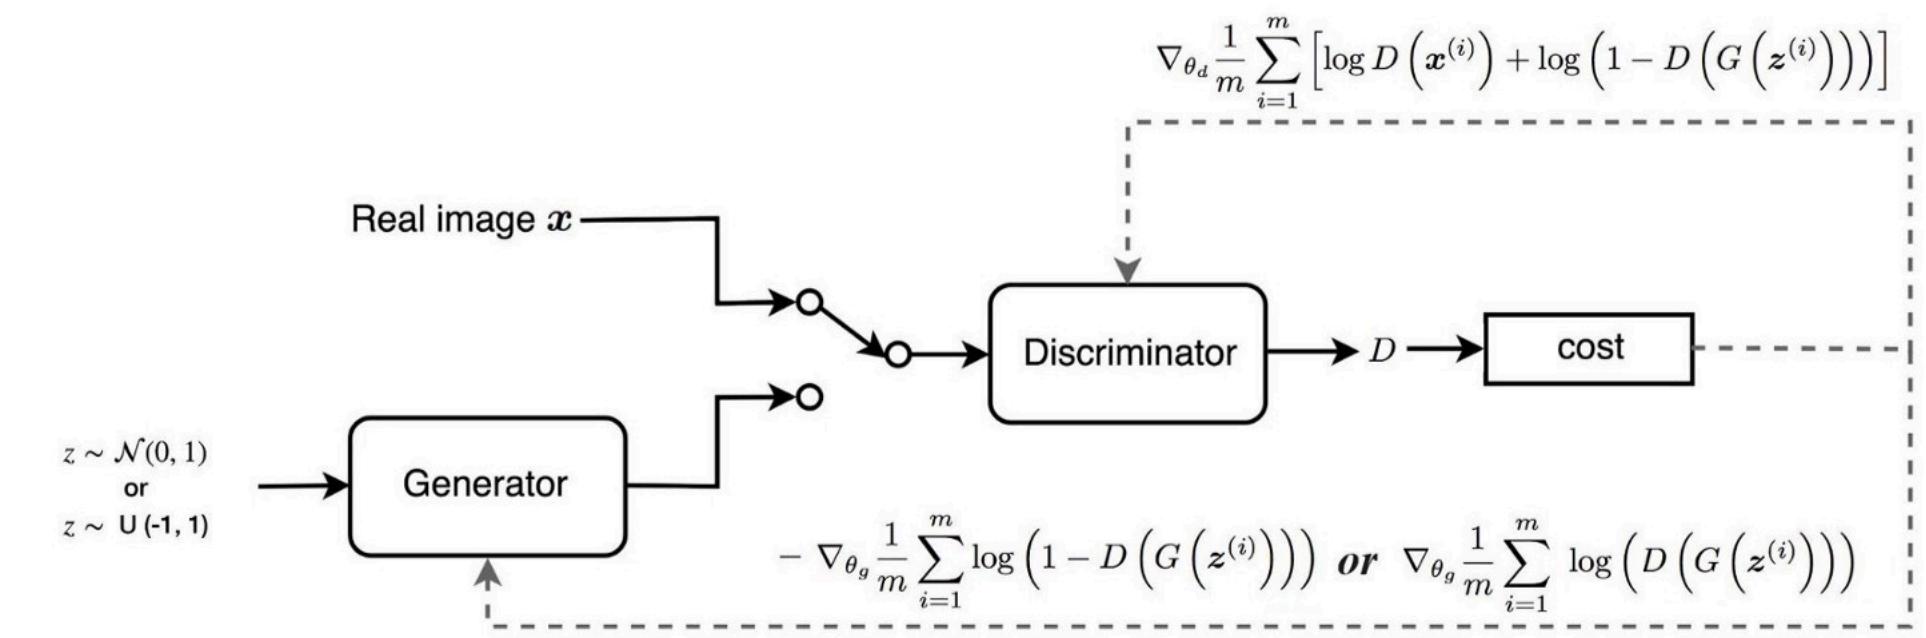
$$\min L = L_c(x, C) + kL_s(x, S)$$

k can be tuned up if you want more style. Gradient descent is

$$\Delta x = -\eta \delta_x L$$



Generative Adversarial Models



General gradient descent algorithm for GAN

Algorithm 1 Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator, k , is a hyperparameter. We used $k = 1$, the least expensive option, in our experiments.

for number of training iterations do

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Sample minibatch of m examples $\{x^{(1)}, \dots, x^{(m)}\}$ from data generating distribution $p_{\text{data}}(x)$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m [\log D(x^{(i)}) + \log (1 - D(G(z^{(i)})))] .$$

end for

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)}))) .$$

end for

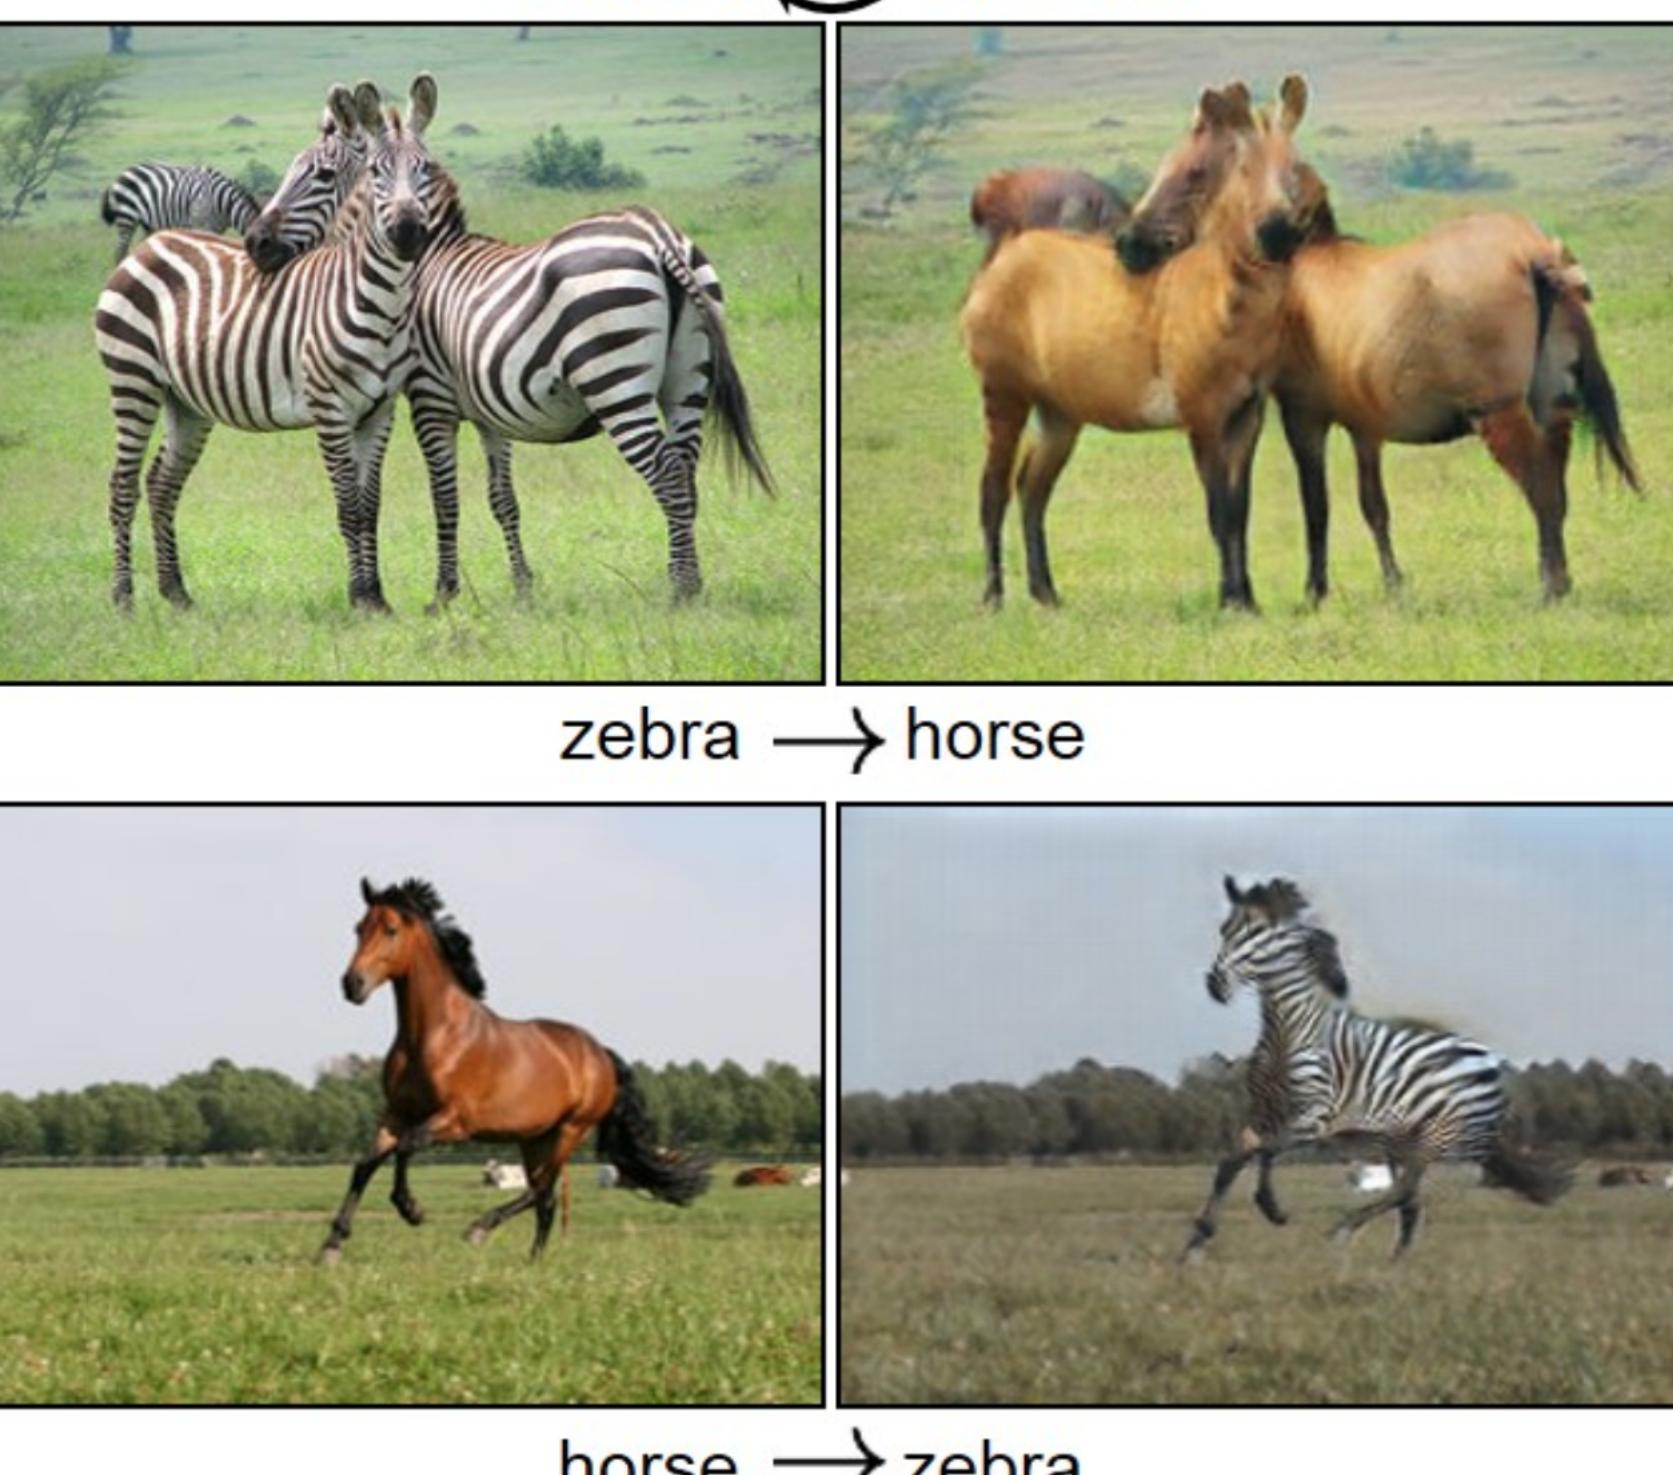
The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

Style-GAN

As the training progresses Style-GAN adds more layers to the generator network G. Initially it starts with one shallow layer to produce low resolution images 4 by 4. After a few iterations of learning 'coarse' features of an image, G adds another layer to produce slightly higher resolution images 8 by 8. It continues doing this, learning finer more detailed features as training progresses until producing HQ 1024 by 1024 images.

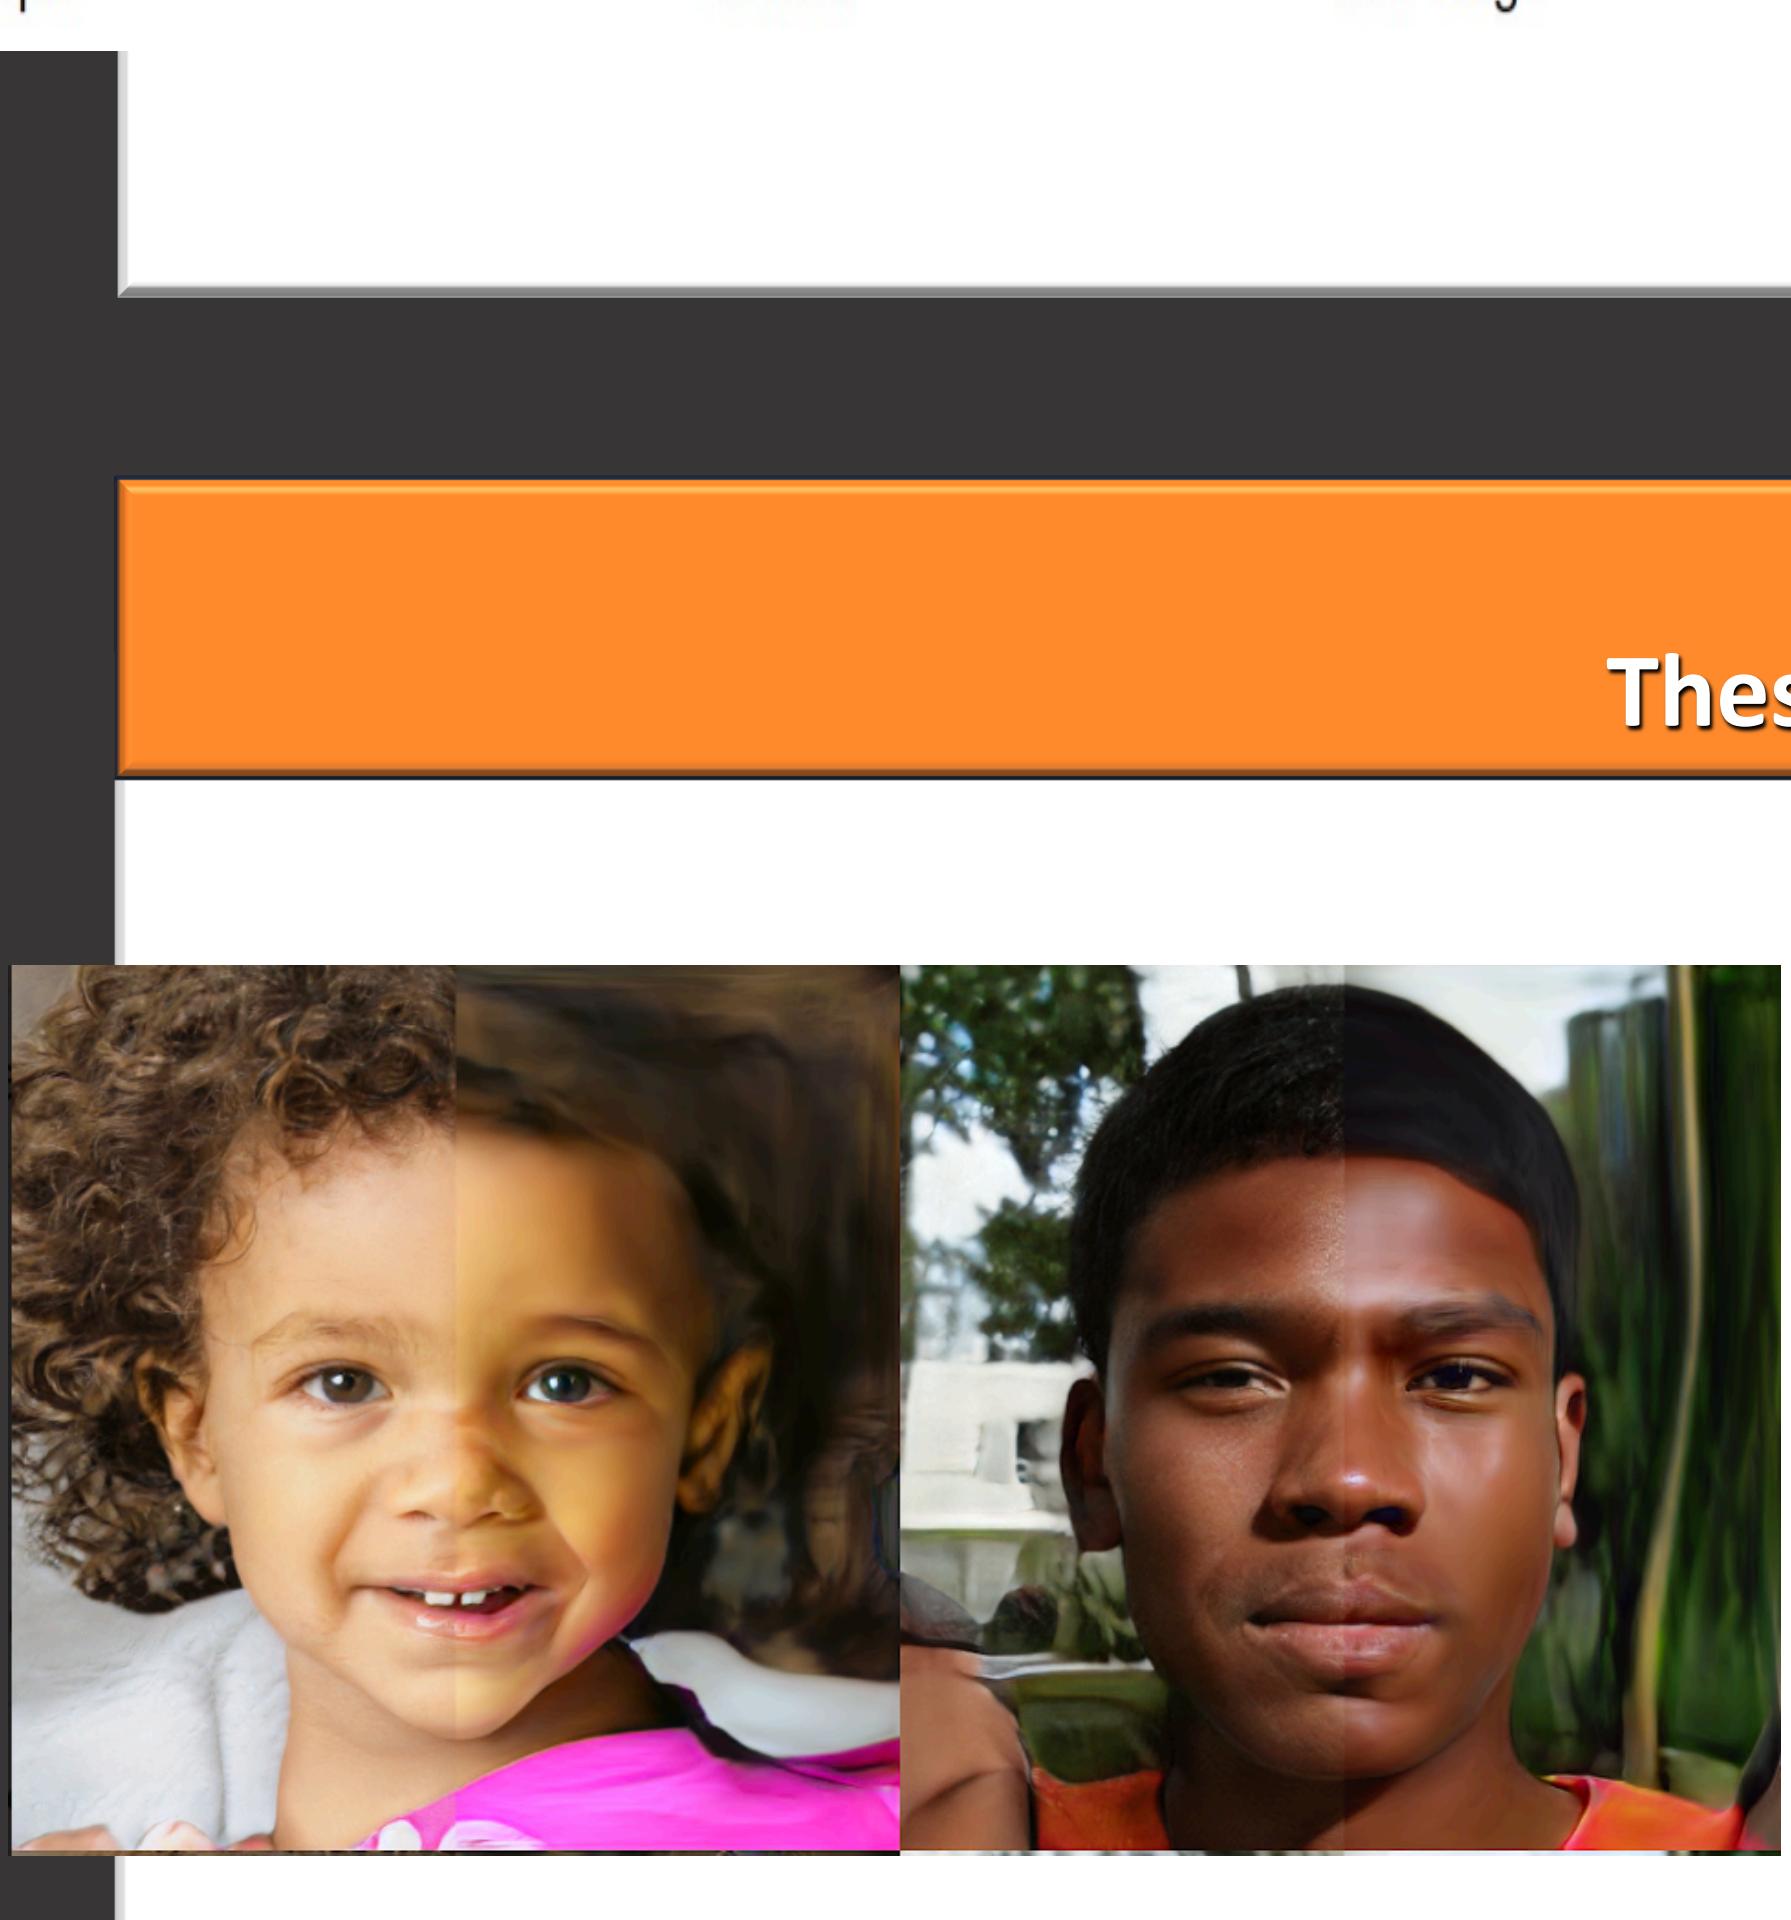


Cycle-GAN



Style-GAN These people (probably) don't exist.

For similar stuff, see <https://thispersondoesnotexist.com/>



References

Leon A. Gatys, Alexander S. Ecker, Matthias Bethge. *A Neural Algorithm of Artistic Style*.

Ian J. Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, et al. *Generative Adversarial Networks*

Jun-Yan Zhu, Taesung Park, Phillip Isola, Alexei A. Efros. UC Berkeley. *Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks*

Tero Karras, Samuli Laine, Timo Aila. *A Style-Based Generator Architecture for Generative Adversarial Networks*

$$\mathcal{L}_{\text{cyc}}(G, F) = \mathbb{E}_{x \sim p_{\text{data}}(x)} [\|F(G(x)) - x\|_1] + \mathbb{E}_{y \sim p_{\text{data}}(y)} [\|G(F(y)) - y\|_1]$$

$$\mathcal{L}_{\text{GAN}}(G, D, X, Y)$$

$$\text{For } G, \text{ minimize } \mathbb{E}_{x \sim p_{\text{data}}(x)} [(D(G(x)) - 1)^2]$$

$$\text{For } D, \text{ minimize } \mathbb{E}_{y \sim p_{\text{data}}(y)} [(D(y) - 1)^2] + \mathbb{E}_{x \sim p_{\text{data}}(x)} [D(G(x))^2]$$

Minimize total loss function: $\mathcal{L}(G, F, D_X, D_Y) = \mathcal{L}_{\text{GAN}}(G, D_Y, X, Y) + \mathcal{L}_{\text{GAN}}(F, D_X, Y, X) + \lambda \mathcal{L}_{\text{cyc}}(G, F)$,